



DataScientest • com

NumPy for Data Science

Introduction to Numpy arrays

Introduction

`numpy` (for *Numerical Python*) is a very advanced digital library for handling large multidimensional arrays and high level mathematical routines (linear algebra, statistics, complex mathematical functions, etc.).

The object class that we will mainly manipulate is the **array** class of `numpy`.

These arrays correspond to N-dimensional matrices which can contain very diverse data such as tabular data, time series or images.

The interest of the `numpy` module lies in the possibility of applying operations on these arrays in a very efficient way, i.e. the number of **lines of code** required and the **computational time** to perform these operations will be very reduced compared to a traditional `Python` syntax.

1. Creating a **numpy** array

Unlike usual classes you will see in your training, a `numpy` array can be instantiated with many different constructors.

The argument of these constructors is usually a **tuple** containing the desired matrix dimensions. This tuple is called the **shape** of an array:

```
# Import of the numpy module under the alias 'np'
import numpy as np

# Creating a matrix of dimensions 5x10 filled with zeros
X = np.zeros(shape = (5, 10))

# Creating a 3-dimensional 3x10x10 matrix filled with ones
X = np.ones(shape = (3, 10, 10))
```

It is also possible to create an array from a **list** using the `np.array` constructor:

```
# Creating an array from a list defined by comprehension
X = np.array([2*i for i in range(10)]) # 0, 2, 4, 6, ..., 18

# Create an array from a list of lists
X = np.array([[1, 3, 3],
              [3, 3, 1],
              [1, 2, 0]])
```

These three constructors are only examples. We will see other constructors later on.

2. Indexation of a **numpy** array

Unlike lists, a `numpy` array is multidimensional. Indexing must be done by entering the index that we want to access **on each dimension**:

```
# Creating a matrix of dimensions 10x10 filled with ones
```

```
X = np.ones(shape = (10, 10))
```


```
# Displaying element at index (4, 3)
```

```
print(X[4, 3])
```

```
# Assigning the value -1 to the element of index (1, 5)
```


```
X[1, 5] = -1
```

As with all other Python indexable objects, the starting index of an axis is 0.



As for lists, it is possible to index an array using **slicing**:





It is possible to slice on **each dimension** of an array. In the following example, we will extract a **subarray** from `X` using slicing:

In [1]:

```
import numpy as np

# Insert your code here
X = np.zeros(shape=(6,6))
X[0:3, 0:3] = 1
X[3:, 3:] = -1
X
```

```
Out[1]: array([[ 1.,  1.,  1.,  0.,  0.,  0.],
               [ 1.,  1.,  1.,  0.,  0.,  0.],
               [ 1.,  1.,  1.,  0.,  0.,  0.],
               [ 0.,  0.,  0., -1., -1., -1.],
               [ 0.,  0.,  0., -1., -1., -1.],
               [ 0.,  0.,  0., -1., -1., -1.]])
```

Hide solution

In [2]:

```
import numpy as np

# Creation of a matrix of shape 6x6 filled with zeros
X = np.zeros(shape = (6, 6))

# Assignment of the value 1 to the first diagonal block
X[:3, :3] = 1

# Assignment of the value -1 to the second diagonal block
X[3:, 3:] = -1

# Displaying the matrix
np.print(X)
```

```
[[ 1.  1.  1.  0.  0.  0.]
 [ 1.  1.  1.  0.  0.  0.]
 [ 1.  1.  1.  0.  0.  0.]
 [ 0.  0.  0. -1. -1. -1.]
 [ 0.  0.  0. -1. -1. -1.]
 [ 0.  0.  0. -1. -1. -1.]]
```

- (b) Using array constructors and slicing, create and display the following matrix:

$$\begin{pmatrix} 0 & 1 & 2 & 3 & 4 & 5 \\ 0 & 1 & 2 & 3 & 4 & 5 \\ 0 & 1 & 2 & 3 & 4 & 5 \\ 0 & 1 & 2 & 3 & 4 & 5 \\ 0 & 1 & 2 & 3 & 4 & 5 \\ 0 & 1 & 2 & 3 & 4 & 5 \end{pmatrix}$$

You can either replace each row of a matrix of zeros with `np.array([0, 1, 2, 3, 4, 5])` or assign to each column its index.

In [19]:

```
# Insert your code here
import numpy as np
x = np.zeros((6,6))
v = [i for i in range(6)]
x[:] = v
x
```

Out[19]: array([[0., 1., 2., 3., 4., 5.],
[0., 1., 2., 3., 4., 5.],
[0., 1., 2., 3., 4., 5.],
[0., 1., 2., 3., 4., 5.],
[0., 1., 2., 3., 4., 5.],
[0., 1., 2., 3., 4., 5.]])

Hide solution

In [17]:

```
# First solution
X = np.zeros(shape = (6, 6))

# We replace each row with 'np.array([0, 1, 2, 3, 4, 5])'
for i in range(6):
    X[i, :] = np.array([0, 1, 2, 3, 4, 5])

# Display the matrix
print("First solution")
print(X)
print("\n")

# Second solution
X = np.zeros(shape = (6, 6))

# Each column of X is assigned its index
for i in range(6):
    X[:, i] = i

# Matrix display
print("Second solution")
print(X)
```

First solution

```
[[0. 1. 2. 3. 4. 5.]
 [0. 1. 2. 3. 4. 5.]
 [0. 1. 2. 3. 4. 5.]
 [0. 1. 2. 3. 4. 5.]
 [0. 1. 2. 3. 4. 5.]
 [0. 1. 2. 3. 4. 5.]]
```

Second solution

```
[[0. 1. 2. 3. 4. 5.]
 [0. 1. 2. 3. 4. 5.]
 [0. 1. 2. 3. 4. 5.]
 [0. 1. 2. 3. 4. 5.]
 [0. 1. 2. 3. 4. 5.]
 [0. 1. 2. 3. 4. 5.]]
```

3. Operations on Numpy arrays: Example

The `numpy` module is not only used to create matrices. Most of the time, the matrices will be created from real data.

The interest of the `numpy` module consists in its optimized **code** which allows to perform computations on large matrices in a very short time.

The `numpy` module contains basic mathematical functions such as:

Function	Numpy Function
e^x	<code>np.exp(x)</code>
<code>log(x)</code>	<code>np.log(x)</code>
<code>sin(x)</code>	<code>np.sin(x)</code>
<code>cos(x)</code>	<code>np.cos(x)</code>
Rounding to <code>n</code> decimals	<code>np.round(x, decimals = n)</code>

The complete list of math operations for `numpy` is given [here \(https://numpy.org/doc/stable/reference/routines.math.html\)](https://numpy.org/doc/stable/reference/routines.math.html).

These functions can be applied to all `numpy` arrays, regardless of their dimensions:

```
X = np.array([i/100 for i in range(100)]) # 0, 0.01, 0.02, 0.03, ..., 0.98, 0.99

# Calculate the exponential of x for x = 0, 0.01, 0.02, 0.03, ..., 0.98, 0.99
exp_X = np.exp(X)
```

In the next cell we will create the array:

In [40]:

```
X = np.array([i/100 for i in range(100)])

# Insert your code here
def f(X):
    result = np.exp(np.sin(X) + np.cos(X))
    return result

print(np.round(f(X)[:10], 2))
```

[2.72 2.75 2.77 2.8 2.83 2.85 2.88 2.91 2.94 2.96]

Hide solution

In [22]:

```
X = np.array([i/100 for i in range(100)])

# Definition of the function f
def f(X):
    return np.exp(np.sin(X) + np.cos(X))

# Calculation of f(X)
result = f(X)

# We round the result to 2 decimal places
rounded = np.round(result, decimals = 2)

# Displaying the first 10 elements of the rounded result
print(rounded[: 10])
```

```
[2.72  2.75  2.77  2.8   2.83  2.85  2.88  2.91  2.94  2.96]
```

- (c) Define a function called `f_python` that performs the same operation on every element of `X` using a `for` loop.

The dimensions of an array `X` can be retrieved using the **shape** attribute of `X` which is a **tuple**: `shape = X.shape` .

For an array with **one** dimension, the number of elements contained in this array corresponds to the **first** element of its shape: `n = X.shape[0]` .

In [36]:

```
# Insert your code here
def f_python(X):
    if len(X.shape) == 1:
        return np.exp(np.sin(X) + np.cos(X))
    else:
        rows, cols = X.shape[0], X.shape[1]
        for r in range(rows):
            for c in range(cols):
                X[r,c] = np.exp(np.sin(X[r,c]) + np.cos(X[r,c]))
    return X

X = np.array([i/100 for i in range(100)])
Y = np.ones((5,4))
Z = np.zeros((6,6))
v = [i for i in range(6)]
Z[:] = v

print(f_python(X[:10]))
print(50*'-')
print(Y)
print(50*'-')
print(f_python(Y))
print(50*'-')
print(f_python(Z))
```

```
[2.71828183 2.74546328 2.7726365  2.79979313 2.8269247  2.85402258
 2.88107805 2.90808226 2.93502626 2.96190097]
```

```
-----
[[1.  1.  1.  1.]
 [1.  1.  1.  1.]
 [1.  1.  1.  1.]
 [1.  1.  1.  1.]
 [1.  1.  1.  1.]]
```

```
-----
[[3.98195654 3.98195654 3.98195654 3.98195654]
 [3.98195654 3.98195654 3.98195654 3.98195654]
 [3.98195654 3.98195654 3.98195654 3.98195654]
 [3.98195654 3.98195654 3.98195654 3.98195654]
 [3.98195654 3.98195654 3.98195654 3.98195654]]
```

```
-----
[[2.71828183 3.98195654 1.63746709 0.42789712 0.24403439 0.50902299]
 [2.71828183 3.98195654 1.63746709 0.42789712 0.24403439 0.50902299]
 [2.71828183 3.98195654 1.63746709 0.42789712 0.24403439 0.50902299]
 [2.71828183 3.98195654 1.63746709 0.42789712 0.24403439 0.50902299]
 [2.71828183 3.98195654 1.63746709 0.42789712 0.24403439 0.50902299]
 [2.71828183 3.98195654 1.63746709 0.42789712 0.24403439 0.50902299]]
```

Hide solution

In [37]:

```
def f_python(X):  
    n = X.shape[0]  
    for i in range(n):  
        X[i] = np.exp(np.sin(X[i]) + np.cos(X[i]))
```

We will now compare the execution times of these two functions applied to a very large array with 10 million values.

We will measure these execution times using the `time` module. To measure the execution time of a function, just take the **difference** between **the start time of execution** and **the end time**. We consider that the assignment of a variable is done instantly.

- (d) Run the next cell to compare execution times. This cell could take some time to run.

In [38]:

```
from time import time  
  
# Creation of an array with 10 million values  
X = np.array([i/1e7 for i in range(int(1e7))])  
  
time_start = time()  
f(X)  
time_end = time()  
  
runtime = time_end - time_start  
  
print ("Calculating f with numpy took", runtime, "seconds")  
  
time_start = time()  
f_python(X)  
time_end = time()  
  
runtime = time_end - time_start  
  
print ("Computing f with a for loop took", runtime, "seconds")
```

Calculating f with numpy took 0.26934242248535156 seconds
Computing f with a for loop took 29.84815502166748 seconds

Conclusion

As you can see, the computation time with a `for` loop is **extremely long**.



Unvalidate