



## Pandas for Data Science

### Introduction to DataFrames

#### Introduction

The `pandas` module has been developed to provide Python with the tools necessary to manipulate and analyze large volumes of data.

Pandas introduces the **DataFrame** class, an array-like data structure that offers more advanced data manipulation and exploration than NumPy arrays.

The main features of `pandas` are:

- data recovery from files (CSV, Excel tables, etc.)
- handling this data (deletion / addition, modification, statistical visualization, etc.).

This notebook aims at:

- Understanding the format of a `DataFrame`.
- Creating a first `DataFrame`.
- Carrying out a first exploration of a dataset using the `DataFrame` class.

- (a) Import the `pandas` module under the name `pd`.

In [1]:

```
# Insert your code here
import pandas as pd
```

Show solution

#### 1. Format of a DataFrame

A DataFrame has an index. Typical identifiers.

A DataFrame is a database (in the sense that features are

The column names are not managed by the user.

The index can be one (or several) columns that we specify.

**Example:** DataFrame with one column and one row.

**Example:** DataFrame with two columns and three rows.

**Example:** DataFrame with one column and three rows.

We will detail the format of a DataFrame in the next section.

The `DataFrame` class has several advantages over a `Numpy` array:

- Visually, a `DataFrame` is much more **readable** thanks to more explicit column and row indexing.
- Within the same column the elements are of the same type but from one column to another, the **type of the elements may vary**, which is not the case of `Numpy` arrays which only support data of the same type.
- The `DataFrame` class contains more methods for handling and preprocessing databases, while `NumPy` specializes instead in optimized computation.

## 2. Creation of a DataFrame: from a NumPy array

It is possible to create a `DataFrame` from a `Numpy` array, which is very practical.

Let's take a concrete example:

```
pd.DataFrame
```

- 
- 
- 

For other information, see the [pandas documentation](https://pandas.pydata.org/pandas-docs/stable/10min.html) (<https://pandas.pydata.org/pandas-docs/stable/10min.html>).

Example:

```
# Create a NumPy array
```

```
# Instantiate a DataFrame from the NumPy array
```

```
# The resulting DataFrame
```

```
']' #
```

This produces the following DataFrame:

## 3. Creation of a DataFrame: from a dictionary

In [3]:

# Insert your code here  
  
dic = {'name' : ['honey', 'flour', 'wine'],  
      'date' : ['08/10/2025', '09/25/2024', '10/15/2023'],  
      'quantity' : [100, 55, 1800],  
      'price' : [2, 3, 10]}  
  
df = pd.DataFrame(dic)  
  
df.head()

Out[3]:

	name	date	quantity	price
0	honey	08/10/2025	100	2
1	flour	09/25/2024	55	3
2	wine	10/15/2023	1800	10

Hide solution

In [ ]:

dictionary = {"Product" : ['honey', 'flour', 'wine'],  
              "Expiration date" : ['10/08/2025', '25/09/2024', '15/09/2024'],  
              "Quantity" : [100, 55, 1800],  
              "Price per unit" : [2, 3, 10]}  
  
df = pd.DataFrame(dictionary)  
  
print(df)

#### 4. Creation of a DataFrame: from a data file

Most often a  
of interest. T

The most cor  
Separated Val  
separated by

Here is an ex

A, B,  
1, 2,  
5, 6,  
9, 10,

In this format

- - 
  -

To import the  
function of p

pd.read  
= 0, i

The essential

- - 
  -

```
In [8]: # Insert your code here

df = pd.read_csv('transactions.csv', sep=',', header=0, index_col=0)
df
```

Out[8]:

	cust_id	tran_date	prod_subcat_code	prod_cat_code	Qty	Rate	Tax	total_a
transaction_id								
80712190438	270351	28-02-2014	1.0	1	-5	-772.0	405.300	-4265.3
29258453508	270384	27-02-2014	5.0	3	-5	-1497.0	785.925	-8270.9
51750724947	273420	24-02-2014	6.0	5	-2	-791.0	166.110	-1748.1
93274880719	271509	24-02-2014	11.0	6	-3	-1363.0	429.345	-4518.3
51750724947	273420	23-02-2014	6.0	5	-2	-791.0	166.110	-1748.1
...	...	...	...	...	...	...	...	...
94340757522	274550	25-01-2011	12.0	5	1	1264.0	132.720	1396.7
89780862956	270022	25-01-2011	4.0	1	1	677.0	71.085	748.0
85115299378	271020	25-01-2011	2.0	6	4	1052.0	441.840	4649.8
72870271171	270911	25-01-2011	11.0	5	3	1142.0	359.730	3785.7
77960931771	271961	25-01-2011	11.0	5	1	447.0	46.935	493.5

23053 rows × 9 columns

Hide solution

```
In [10]: # You can directly specify the name of the column containing the index

transactions = pd.read_csv(filepath_or_buffer = 'transactions.csv',
                             sep = ',',
                             header = 0,
                             index_col = 'transaction_id')

# You can also directly enter the number of the column that indexes the data

transactions = pd.read_csv(filepath_or_buffer = 'transactions.csv',
                             sep = ',',
                             header = 0,
                             index_col = 0) # number of the column that indexes the data
```

We loaded the `transactions.csv` file in the DataFrame `transactions` which gathers a history of transactions carried out between 2011 and 2014. In the next section, we will study this dataset.

### 5. First exploration of a dataset using the `DataFrame` class

The rest of the DataFrame that is:

- 
- 
- 

As a reminder, a DataFrame is a 2D array-like object. Example:

### 6. Visualization of the data using columns and rows

- It is possible to select lines of the DataFrame.

For that, we use the `loc` method. The `loc` method returns a new DataFrame with the selected lines.

It is also possible to select columns of the DataFrame. The `loc` method is applied in the following example:

```
# Display the first 5 lines of the DataFrame
my_data = transactions.head(5)
```

In [50]:

```
# Insert your code here
transactions.head()
```

Out[50]:

	cust_id
transaction_id	
80712190438	270351
29258453508	270384

Show solution

- (b) Display the last 5 lines of the DataFrame.

In [51]:

```
# Insert your code here
transactions.tail(1)
```

Out[51]:

	cust_id	tran_date	prod_subcat_code	prod_cat_code	Qty	Rate	Tax	total_amt
transaction_id								
77960931771	271961	25-01-2011		11.0	5	1	447.0	46.935

Hide solution

In [ ]:

```
transactions.tail(10)
```

It is possible to retrieve the **name of the columns** of a `DataFrame` thanks to its **columns** attribute.

```
# Creation of a DataFrame df from a dictionary
dictionary = {'A': [1, 5, 9],
              'B': [2, 6, 10],
              'C': [3, 7, 11],
              'D': [4, 8, 12]}
```

```
df = pd.DataFrame (data = dictionary, index = ['i_1', 'i_2', 'i_3'])
```

These instructions produce the same `DataFrame` as before:

	A	B	C	D
i_1	1	2	3	4
i_2	5	6	7	8
i_3	9	10	11	12

```
# Display of df DataFrame columns
print(df.columns)
>>> ['A', 'B', 'C', 'D']
```

The list of the column names can be used to iterate over the columns of a `DataFrame` within a loop.

It can be interesting to know how many transactions (rows) and how many features (columns) the dataset contains.

For this we will use the **shape** attribute of the `DataFrame` class which displays the **dimensions** of our `DataFrame` in the form of a tuple (number of rows, number of columns):

```
# Display the dimensions of df
print (df.shape)
>>> (3,4)
```

- (c) Display the dimensions of the dataset (number of rows, number of columns)

In [15]:

```
# Insert your code here
print(transactions.tail(10))
```

(23053, 9)

Out[15]: 'Qty'

Hide solution

In [16]:

```
print(transactions.tail(10))
print(transactions.columns)
```

(23053, 9)  
Qty

## 7. Selecting columns

In [23]:

```
### Insert your code here

cat_vars = transactions[['cat_id']]
num_vars = transactions[['num_id']]

print(cat_vars.head(3))
print(30*'-')
print(num_vars.head(3))
```

Extracting columns from a `DataFrame` is almost identical to extracting data from a dictionary.

To extract a **column** from a `DataFrame`, all we have to do is enter **between brackets** the **name** of the column to extract. To extract **several** columns, we must enter between brackets **the list of the names** of the columns to extract:

```
# Display of the 'cust_id' column
print(transactions['cust_id'])

# Extraction of 'cust_id' and 'Qty' columns from transactions
cust_id_qty = transactions[["cust_id", "Qty"]]
```

`cust_id_qty` is a new **DataFrame** containing only the `'cust_id'` and `'Qty'` columns.

The display of the first 3 lines of `cust_id_qty` yields:

	cust_id	Qty
transactions_id		
80712190438	270351	-5
29258453508	270384	-5
51750724947	273420	-2

```
\
transaction_id
80712190438
29258453508
51750724947
93274880719
51750724947
```

```
transaction_id
80712190438
29258453508
51750724947
93274880719
51750724947
-----
```

```
transaction_id
80712190438
29258453508
51750724947
93274880719
51750724947
```

Hide solution

In [24]:

```
# Extraction of categorical variables
cat_var_names = ['cust_id', 'tran_date', 'prod_subcat_code', 'prod_cat_code']
cat_vars = transactions[cat_var_names]

# Extraction of quantitative variables
num_var_names = ['Qty', 'Rate', 'Tax', 'total_amt']
num_vars = transactions[num_var_names]

# Display of the first 5 lines of each DataFrame
print ("Categorical variables: \n")
print (cat_vars.head(), "\n \n")

print ("Quantitative variables: \n")
print (num_vars.head())
```

Categorical variables:

	cust_id	tran_date	prod_subcat_code	prod_cat_code
\				
transaction_id				
80712190438	270351	28-02-2014	1.0	1
29258453508	270384	27-02-2014	5.0	3
51750724947	273420	24-02-2014	6.0	5
93274880719	271509	24-02-2014	11.0	6
51750724947	273420	23-02-2014	6.0	5

	Store_type
transaction_id	
80712190438	e-Shop
29258453508	e-Shop
51750724947	TeleShop
93274880719	e-Shop
51750724947	TeleShop

Quantitative variables:

	Qty	Rate	Tax	total_amt
transaction_id				
80712190438	-5	-772.0	405.300	-4265.300
29258453508	-5	-1497.0	785.925	-8270.925
51750724947	-2	-791.0	166.110	-1748.110
93274880719	-3	-1363.0	429.345	-4518.345
51750724947	-2	-791.0	166.110	-1748.110

## 8. Selecting rows of a DataFrame: loc and iloc methods

To extract on  
loc is a ver  
between squ  
very similar t

In order to re  
enter i as a

```
# We r  
e num_  
print(  

```

```
>>  
>> tra  
>> 807  
>> 807
```

In order to re

```
# We r  
258453  
ataFra  
transa  
724947
```

loc can als  
refine the dat

```
# We e  
om the  
transa  
x', 't
```

This instructi

The iloc m  
array, that is  
columns. This

s instruction produces the following DataFrame :

If the row indexing is the one by default (row numbering), the `loc` and `iloc` methods are **equivalent**.

As with Numpy arrays, we can use **conditional indexing** to extract rows from a `Dataframe` that meet a given condition.

```
df[df.loc[df['col 2'] == 3]
```

There are two syntaxes for conditionally indexing a `DataFrame` :

```
df.loc[df['col 2'] == 3]
```

Indeed, indexing with the syntax `df[df['col_2'] == 3]` only returns a **copy** of these entries and does not provide access the memory location where the data is located.

In [26]:

Out[26]:

Hide solution

In [ ]:

Now, the manager would like to find the employee whose identifier is 20.

- (d) In a DataFrame with client identifier
- (e) A column in a value in df[ compute and display



In [30]:

```
# Insert your code here

transactions_client_268819 = transactions.loc[transactions.cust_id == 268819]

# First Way
total = 0
for i in transactions_client_268819['total_amt']:
    total += i
print(total)

# Second Way
transactions_client_268819['total_amt'].sum()
```

14911.974999999999

Out[30]: 14911.974999999999

Hide solution

In [29]:

```
# Extraction of the transactions of the customer which identifier is 268819
transactions_client_268819 = transactions[transactions['cust_id'] == 268819]

# Computation of the total amount of transactions
total = 0

# For each amount in the column 'total_amt'
for amount in transactions_client_268819['total_amt']:
    # We sum the amounts
    total += amount

print(total)
```

14911.974999999999

## 10. Quick statistical study of the data in a DataFrame.

The **describe** method of a `DataFrame` returns a summary of the **descriptive statistics** (min, max, mean, quantiles,...) of its **quantitative** variables. It is therefore a very useful tool for a first visualisation of the type and distribution of these variables.

To analyse the **categorical** variables, it is recommended to start by using the **value\_counts** method which returns the **number of occurrences** for each modality of these variables. The `value_counts` method cannot be used directly on a `DataFrame` but only on the columns of the `DataFrame` which are objects of the **pd.Series** class.

- (a) Use the `describe` method of the `DataFrame` `transactions`.
- (b) The quantitative variables of `transactions` are 'Qty', 'Rate', 'Tax' and 'total\_amt'. By default, are the statistics produced by the `describe` method **only** computed on the quantitative variables?
- (c) Display the number of occurrences of each modality of the `Store_type` column

using the `value_counts` method

In [32]:

```
# Insert your code here

# a
transactions.describe()

# b
# No some categorical variables
# categorical variables
```

Out[32]:

	cust_id
count	23053.000000
mean	271022.241661
std	2430.830508
min	266783.000000
25%	268936.000000
50%	270981.000000
75%	273114.000000
max	275265.000000

In [33]:

```
# c)
transactions.Store_type.value_counts()
```

Out[33]:

e-Shop  
MBR  
Flagship store  
TeleShop  
Name: Store\_type, dtype: object

Hide solution

In [ ]:

```
transactions.describe()

transactions['Store_type'].value_counts()
```

The `describe` method computed statistics on the variables `cust_id`, `prod_subcat_code` and `prod_cat_code` while these are **categorical** variables.

Of course, these statistics make **no sense**. The `describe` method has treated these variables as quantitative because the modalities they take are

In [37]:

```
# Insert your code here

print('average total amount spent :', transactions['total_amt'].mean)
print('maximum quantity purchased :', transactions['Qty'].max())

transactions.describe()
```

average total amount spent : 2108.007266944553  
maximum quantity purchased : 5

Out[37]:

	cust_id	prod_subcat_code	prod_cat_code	Qty	Rate	Tax
count	23053.000000	23021.000000	23053.000000	23053.000000	23031.000000	23031.000000
mean	271022.241661	6.150298	3.762721	2.434173	636.405019	248.665526
std	2430.830508	3.726557	1.677314	2.265703	622.053592	187.087709
min	266783.000000	1.000000	1.000000	-5.000000	-1499.000000	7.350000
25%	268936.000000	3.000000	2.000000	1.000000	312.000000	98.175000
50%	270981.000000	5.000000	4.000000	3.000000	710.000000	199.080000
75%	273114.000000	10.000000	5.000000	4.000000	1109.000000	365.820000
max	275265.000000	12.000000	6.000000	5.000000	1500.000000	787.500000

Hide solution

In [ ]:

```
# Applying the describe method to the transactions DataFrame
transactions.describe()

# The average total amount spent is €2109.
# The maximum quantity purchased is 5.
```

Some transactions have **negative** amounts.

These are transactions that have been cancelled and refunded to the client. These amounts will disrupt the distribution of the amounts which gives us **bad estimates** of the mean and quantiles of the variable `total_amt`.

- (f) What is the average amount of transactions with **positive** amounts?

In [48]:

```
# Insert your code here

print('Average amount of transactions')
```

Average amount of transactions: 6676

Hide solution

In [45]:

```
transactions['total_amt'].mean()

# the average amount spent is €2109
# €500 more than the average
```

Out[45]:

	cust_id
count	20861.000000
mean	271027.632760
std	2432.810074
min	266783.000000
25%	268938.000000
50%	271009.000000
75%	273122.000000
max	275265.000000

Conclusion and

The `DataFrame` class of the `pandas` module will be your favorite data structure when exploring, analysing and processing datasets and databases.

In this brief introduction, you have learned to:

- Create a `DataFrame` from a `numpy` array and a dictionary using the **`pd.DataFrame`** constructor.
- Create a `DataFrame` from a `.csv` file using the **`pd.read_csv`** function.



Validate